

Tensorflow 工具介绍

Introduction to Tensorflow

Ivan

Aug. 27th, 2017

法律声明

本课件、大数据文摘x犀牛学院网站、社群内所有内容，包括但不限于演示文稿、软件、声音、图片、视频、代码等，由发布者本人和大数据文摘共同享有。未经大数据文摘明确书面授权，任何人不得翻录、发行、播送、转载、复制、重制、改动或利用大数据文摘的局部或全部内容或服务，不得在非大数据文摘所属的服务器上作镜像，否则以侵权论，依法追究法律责任。本网站享有对用户在本网站活动的监督和指导权，对从事非法活动的用户，有权终止对其所有服务。



课程合作请联系

info@bigdatadigest.cn

Disclaimer: Slides based on Stanford CS20, Tensorflow for Deep Learning Research

扫码添加客服进群交流

微信15510583366

Outline

- Overview of TensorFlow
- Graphs and Sessions
- TensorBoard Introduction
- Constants and Placeholders
- MNIST Classification with logistic regression using TensorFlow

TensorFlow

- What's TensorFlow?
 - Open source library for numeric and symbolic computation with GPU support
 - Developed by the Google Brain Team for machine learning related research

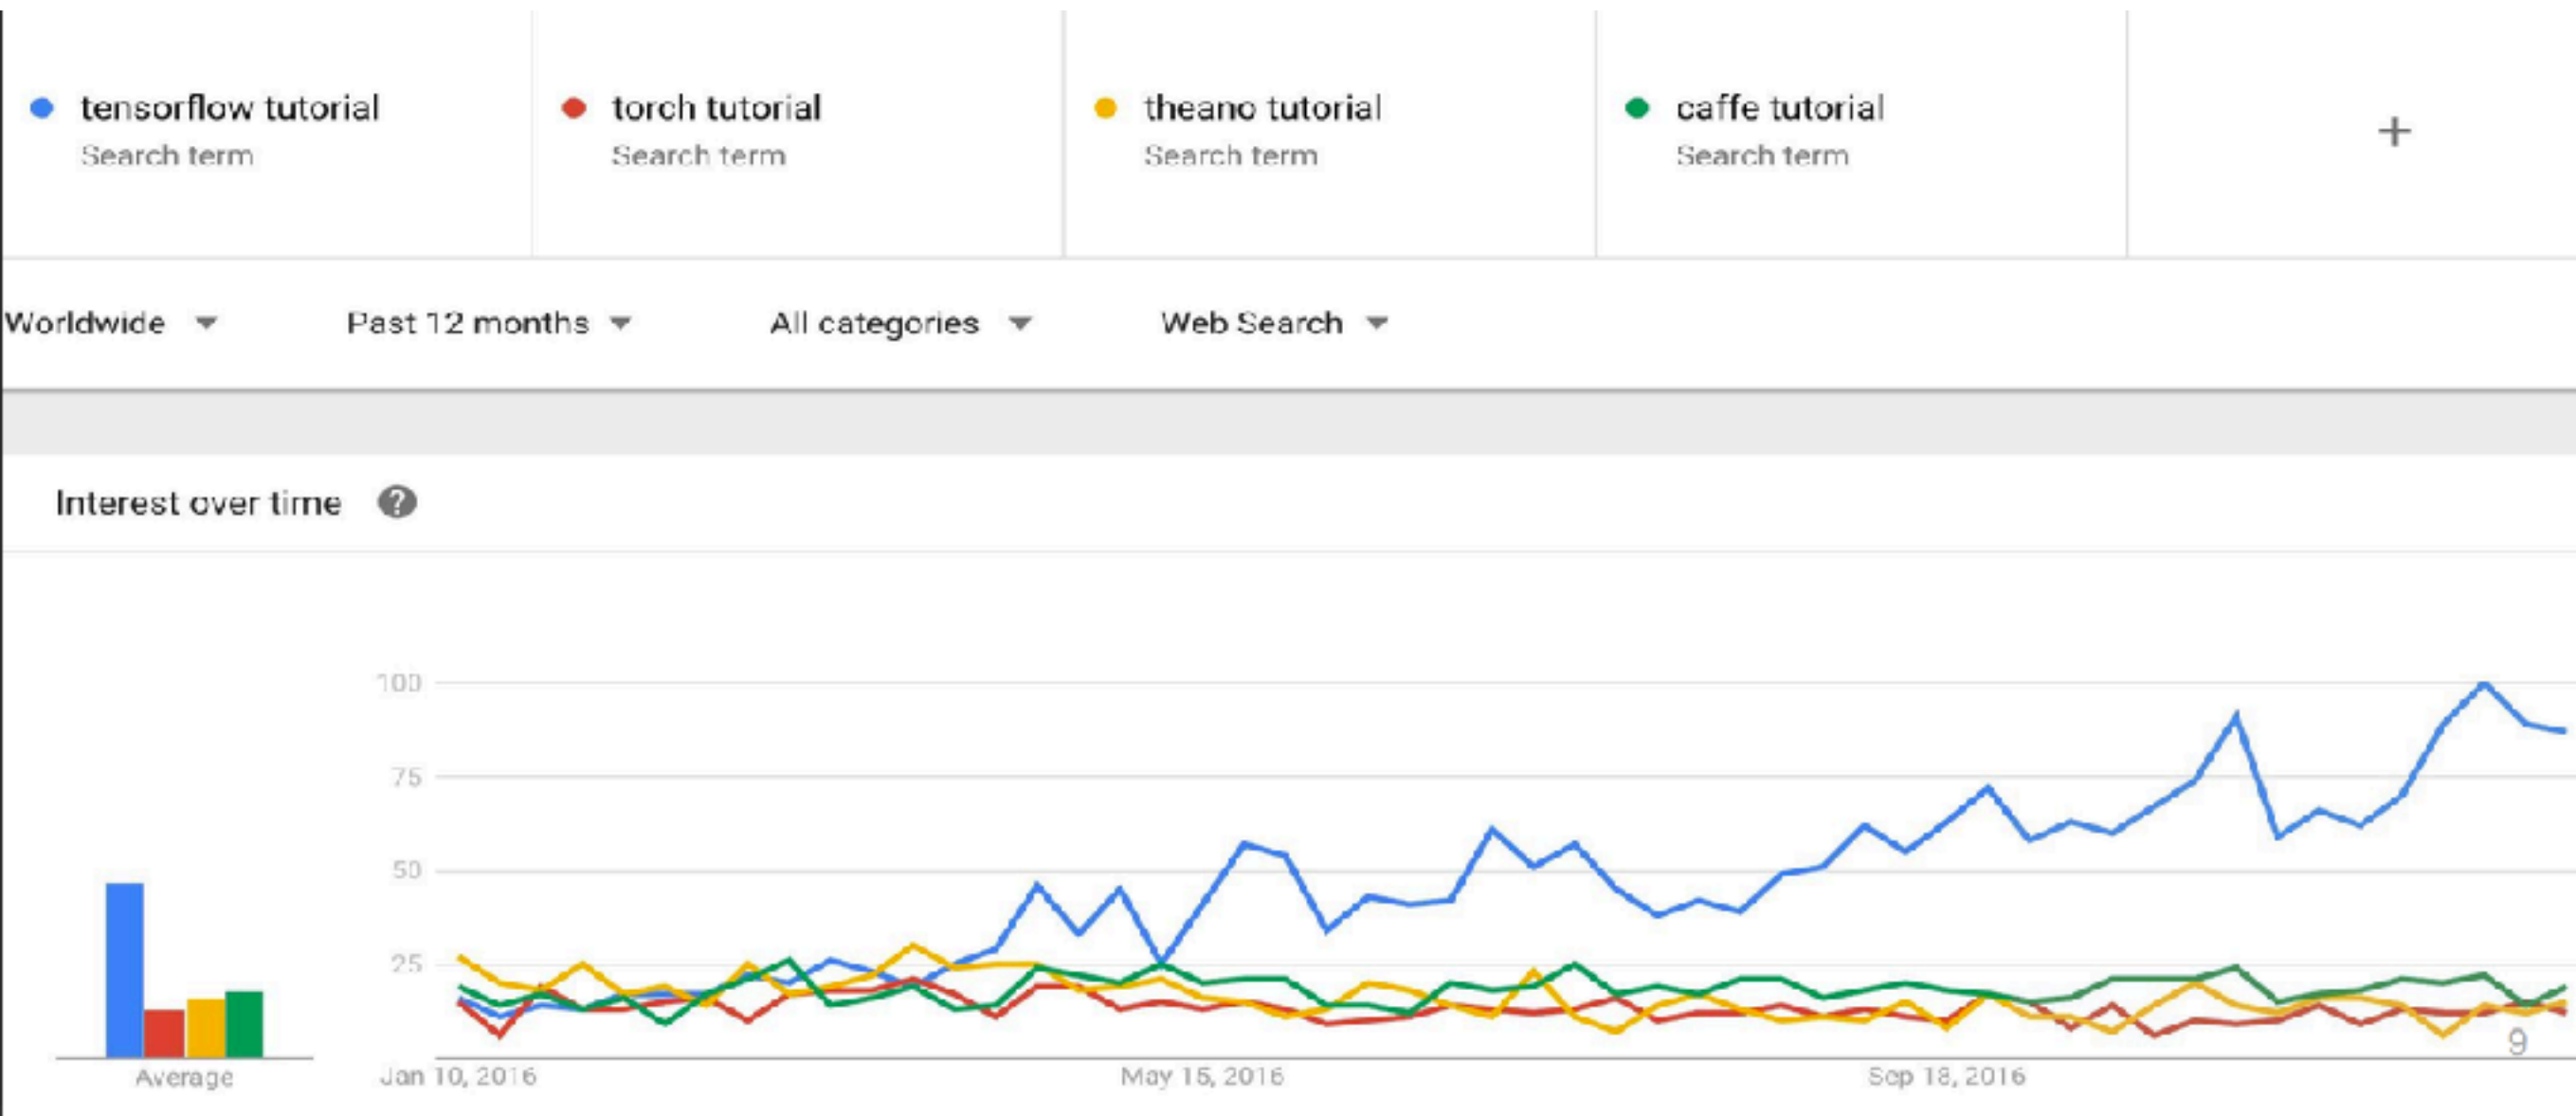
We will compare TensorFlow with Numpy to highlight its key features

TensorFlow

- Besides TensorFlow
- Caffe / Caffe 2 (Berkeley / Facebook)
- Torch (NYU)
- Theano (University of Montreal)
- CNTK (Microsoft)
- Paddle (Baidu)
- MXNet (Amazon)
- PyTorch (Facebook)

TensorFlow

- Why TensorFlow?



TensorFlow

- Why TensorFlow?
- Python API
- Portability: easy to deploy computations over one or more CPUs/GPUs, with the same API
- Flexibility: easy to extend to mobile devices, including Android, iOS, etc.
- Visualization: TensorBoard is great!
- **Auto-differentiation: autodiff, no need to compute the gradient manually**
- Large community: > 10,000 commits and > 3,000 TF-related repos in 1 year

TensorFlow

- Companies that use TensorFlow
 - Google
 - DeepMind
 - Dropbox
 - Snapchat
 - Uber
 - eBay
 - OpenAI
 - ...

TensorFlow

- Goals of this lecture
- Understand TF's computation graph approach
- Explore TF's built-in functions
- Be familiar with the pipeline of a typical machine learning project (MNIST image classification using TF)

Introduction


```
import tensorflow as tf
```

Introduction

- Even higher level abstraction of TF:
- TF Learn (tf.contrib.learn): simplified interface of TensorFlow, similar to scikit-learn
- TF Slim (tf.contrib.slim): lightweight library for defining and running complex models in TensorFlow
- Even more: **Keras**

Graphs and Sessions


- Data flow graphs



Defining computation \neq Execution of Computation

Graphs and Sessions

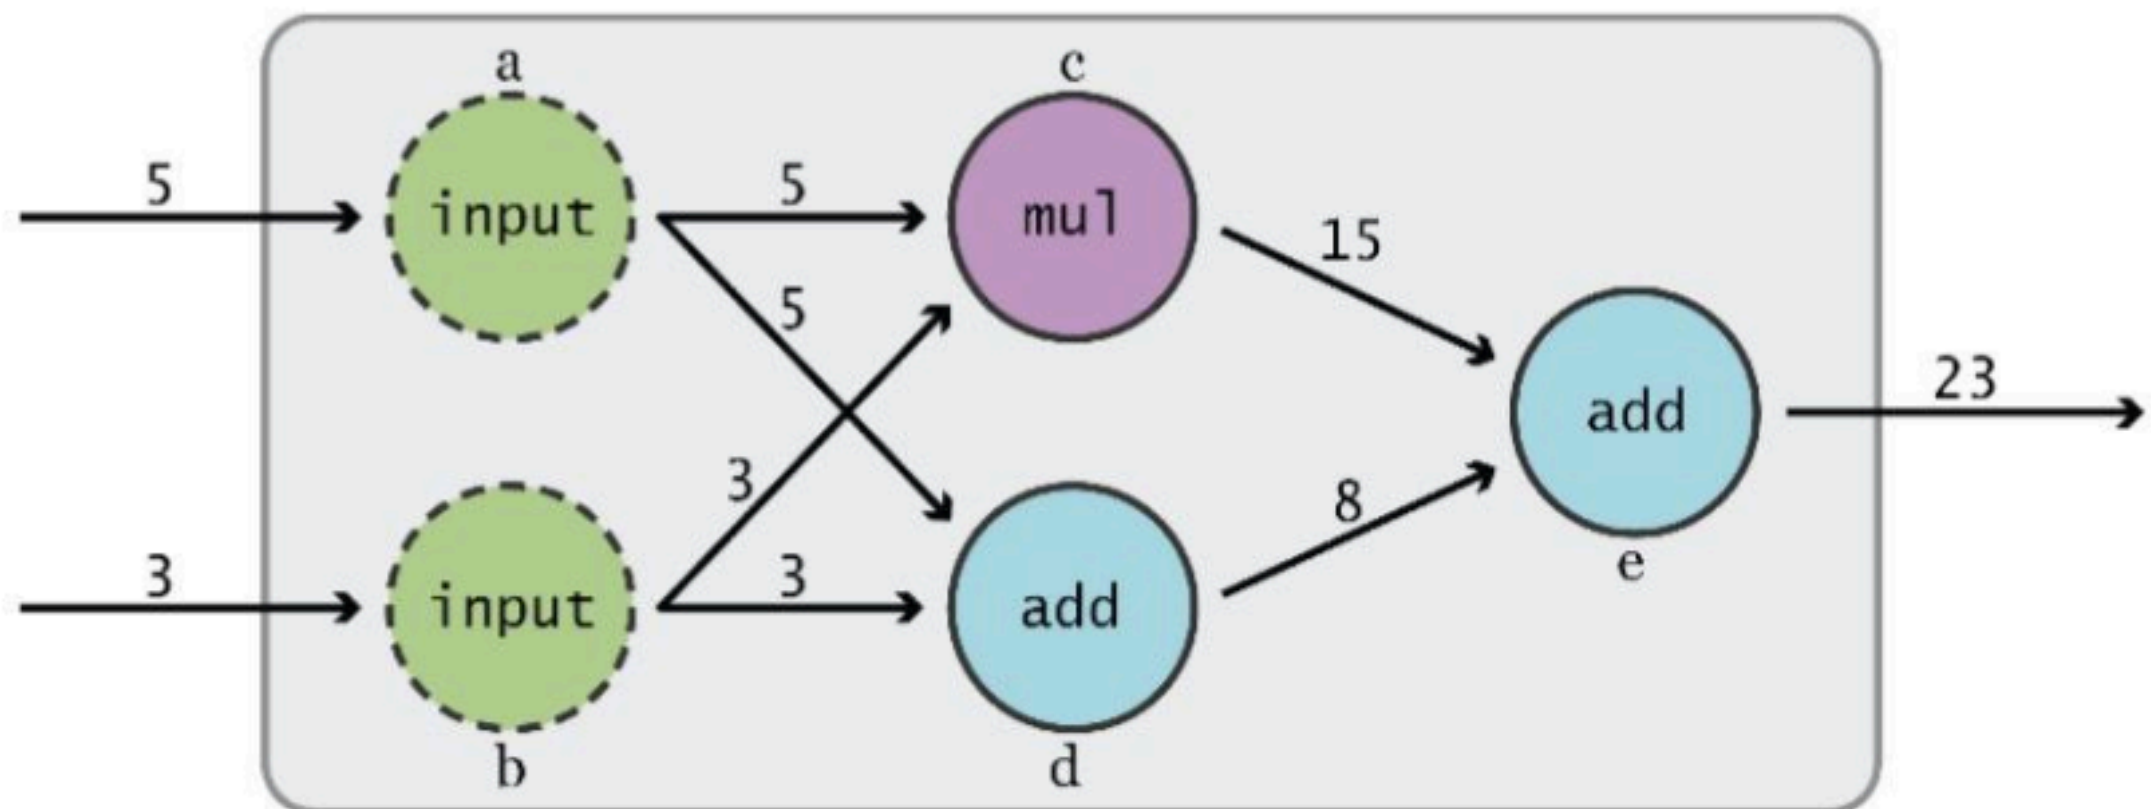
- Data flow graphs



- Typical pipeline in TF:
 - Define the computation graph
 - Run session to execute the computational graph

Graphs and Sessions

- Data flow graphs



- TensorFlow = Tensor + Flow:
 - Tensor = data
 - Flow = operators

Data are transmitted and transformed by operators (op) in the computational graph

TensorFlow

- What is a tensor?
- An n-dimensional array
 - 0-d tensor: scalar (number)
 - 1-d tensor: vector
 - 2-d tensor: matrix
 - ...

TensorFlow

- Numpy vs TensorFlow

Numpy	TensorFlow
<code>a = np.zeros((2,2)); b = np.ones((2,2))</code>	<code>a = tf.zeros((2,2)); b = tf.ones((2,2))</code>
<code>np.sum(b, axis=1)</code>	<code>tf.reduce_sum(b, reduction_indices=[1])</code>
<code>a.shape</code>	<code>a.get_shape()</code>
<code>np.reshape(a, (1, 4))</code>	<code>tf.reshape(a, (1, 4))</code>
<code>5*b+1</code>	<code>5*b+1</code>
<code>np.dot(a, b)</code>	<code>tf.matmul(a, b)</code>
<code>a[1, 1], a[:, 1], a[1, :]</code>	<code>a[1, 1], a[:, 1], a[1, :]</code>

TensorFlow


- Data flow graphs

```
import tensorflow as tf  
a = tf.add(3, 5)
```

What is x, y?

$x = 3, y = 5;$

TF will automatically name variables



Visualization from TensorBoard


Nodes: operators, variables, or constants

Edges: tensors

TensorFlow

- Data flow graphs

```
import tensorflow as tf  
a = tf.add(3, 5)  
print a
```




Visualization from TensorBoard

```
>> Tensor("Add:0", shape=(), dtype=int32)  
(Not 8)  
Why?
```

TensorFlow

- Data flow graphs

```
import tensorflow as tf  
a = tf.add(3, 5)  
print a
```



Visualization from TensorBoard


```
>> Tensor("Add:0", shape=(), dtype=int32)  
(Not 8)
```

Symbolic variable! How to get the value of a?

TensorFlow

- Data flow graphs

```
import tensorflow as tf  
a = tf.add(3, 5)  
print a
```




Visualization from TensorBoard

We need to create a **session** in order to get the value of a

TensorFlow

- Create a session

```
import tensorflow as tf  
a = tf.add(3, 5)  
sess = tf.Session()  
print sess.run(a)    >> output 8  
sess.close()
```




Session will find the dependency of `a`, and computes all the nodes that lead to `a`

TensorFlow

- Create a session (Recommended practice)

```
import tensorflow as tf
a = tf.add(3, 5)
sess = tf.Session()
print sess.run(a)
sess.close()
```



```
import tensorflow as tf
a = tf.add(3, 5)
with tf.Session() as sess:
    print sess.run(a)
```

Session will find the dependency of a, and computes all the nodes that lead to a

TensorFlow

- Summary

A Session object encapsulates the running environment such that operators are executed and tensors are evaluated

TensorFlow

- More examples

```
import tensorflow as tf
```

```
x = 2
```

```
y = 3
```

```
op1 = tf.add(x, y)
```


```
op2 = tf.multiply(x, y)
```

```
op3 = tf.pow(op2, op1)
```

```
with tf.Session() as sess:
```

```
    op3 = sess.run(op3)
```

```
print op3    >> output 7776
```



TensorFlow

- More examples

```
import tensorflow as tf
```

```
x = 2
```

```
y = 3
```

```
add_op = tf.add(x, y)
```

```
mul_op = tf.multiply(x, y)
```


```
useless = tf.multiply(x, add_op)
```

```
pow_op = tf.pow(add_op, mul_op)
```

```
with tf.Session() as sess:
```

```
    z = sess.run(pow_op)
```

```
print z    >> output 15625
```



Think: will session also compute the value of useless?

TensorFlow

- More examples

```
import tensorflow as tf
```

```
x = 2
```

```
y = 3
```

```
add_op = tf.add(x, y)
```

```
mul_op = tf.multiply(x, y)
```

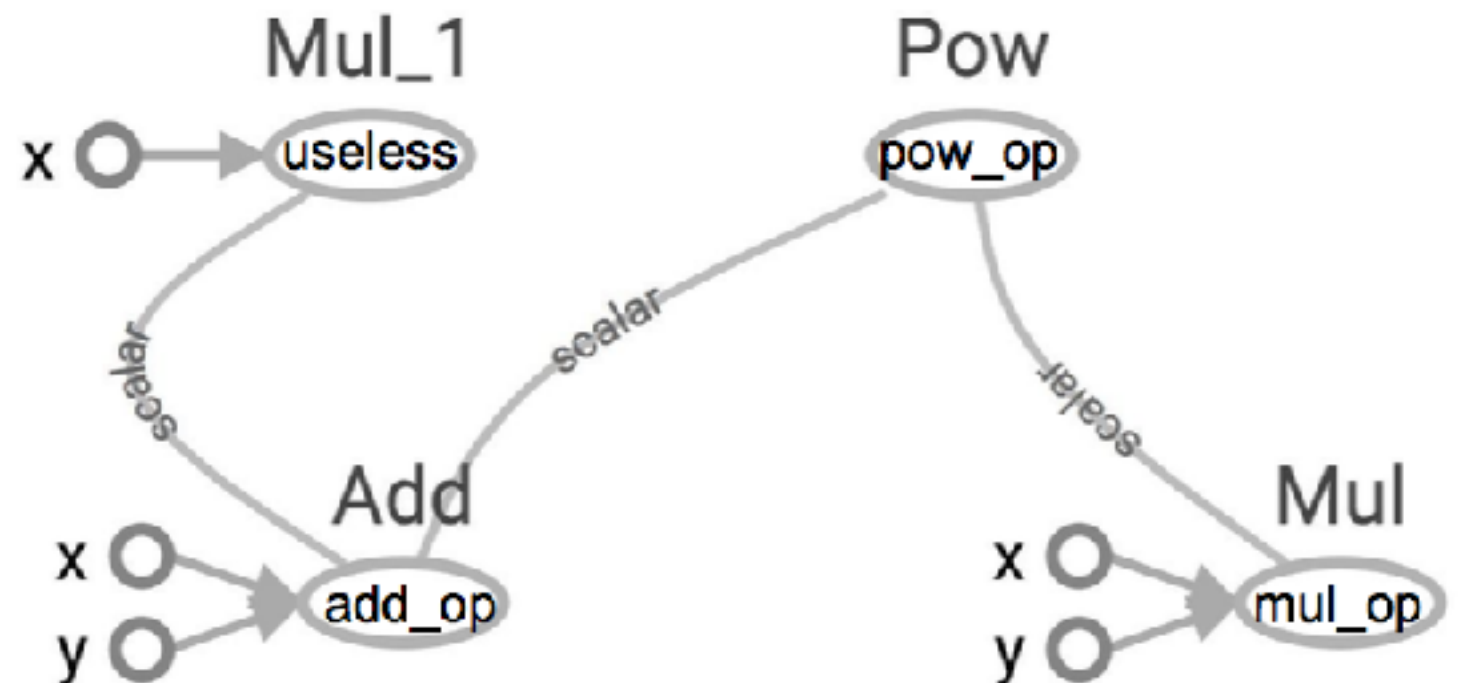
```
useless = tf.multiply(x, add_op)
```

```
pow_op = tf.pow(add_op, mul_op)
```

```
with tf.Session() as sess:
```

```
    z = sess.run(pow_op)
```

```
print z    >> output 15625
```



No: `print useless >> Tensor("Mul_3:0", shape=(), dtype=int32)`

TensorFlow

- More examples

```
import tensorflow as tf
```

```
x = 2
```

```
y = 3
```

```
add_op = tf.add(x, y)
```

```
mul_op = tf.multiply(x, y)
```

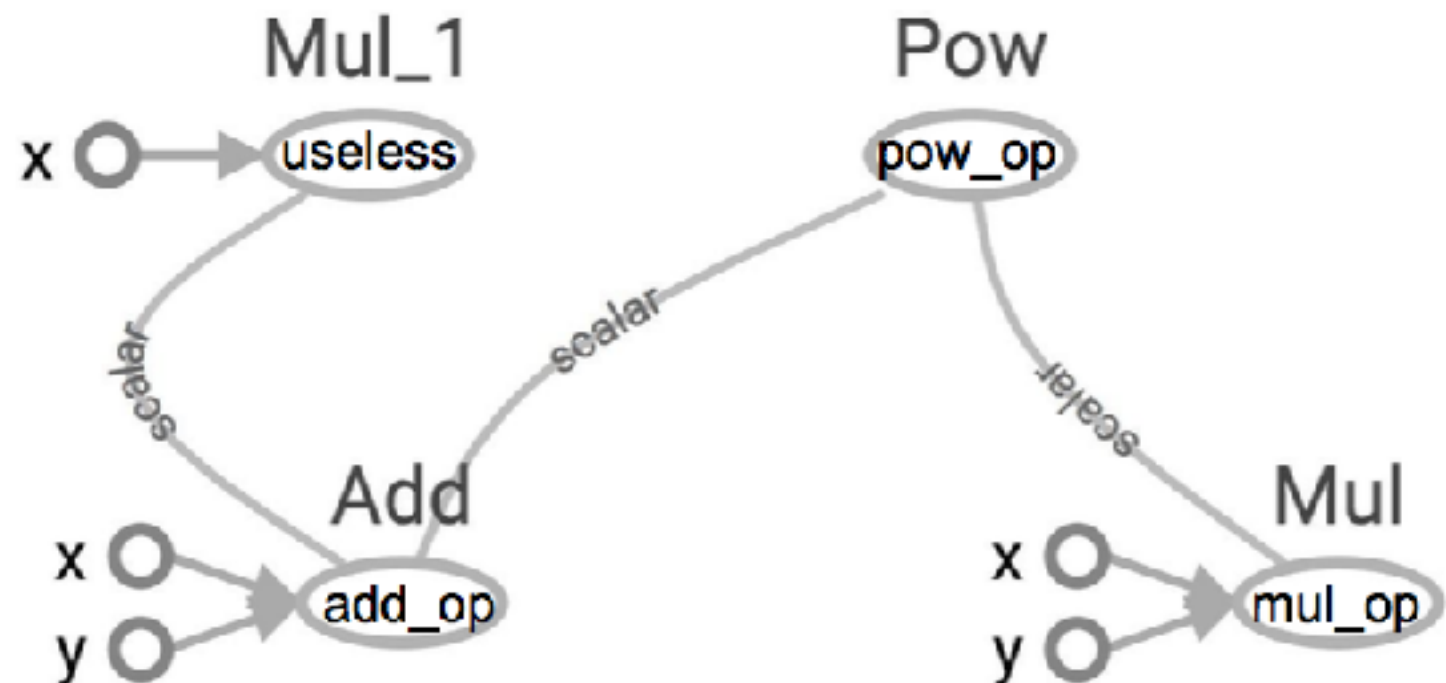
```
useless = tf.multiply(x, add_op)
```

```
pow_op = tf.pow(add_op, mul_op)
```

```
with tf.Session() as sess:
```

```
    z, w = sess.run([pow_op, useless])
```

```
print z, w    >> output 15625, 10
```



API: `tf.Session.run(fetches, feed_dict=None, options=None, run_metadata=None)`

Pass all the variables you want to evaluate to a list in `fetches`

TensorFlow

- Run session with specific device:

```
import tensorflow as tf
# build computational graph.
with tf.device("/gpu:0"):
    a = tf.constant([1.0, 2.0, 3.0, 4.0], shape=(2, 2), name="a")
    b = tf.constant([2.0, 4.0, 6.0, 8.0], shape=(2, 2), name="b")
    c = tf.matmul(a, b)

# build session, set log_device_placement=True
with tf.Session(config=tf.ConfigProto(allow_soft_placement=True)):
    print sess.run(c)
```

TensorFlow

- Why computational graph?
- Save computations (only evaluates subgraphs which lead to the values you're interested in)
- Facilitate distributed computation (model parallelization)
- Directed acyclic graph is required in order to implement auto-differentiation

TensorFlow

- TensorBoard for visualization

```
import tensorflow as tf
a = tf.constant(2)
b = tf.constant(3)
x = tf.add(a, b)
with tf.Session() as sess:
    # add this line to use TensorBoard
    writer = tf.summary.FileWriter("./graphs", sess.graph)
    print sess.run(x)
writer.close()
```

Create the summary writer after graph definition but before running the session

TensorFlow

- TensorBoard for visualization

Open terminal, run:

```
$ python [thisprogram].py
```

```
$ tensor board --logdir="./graphs" --port 5555
```

TensorFlow

The screenshot shows the TensorBoard web interface running on a browser at localhost:5555. The interface has an orange top navigation bar with tabs for SCALARS, IMAGES, AUDIO, GRAPHS (selected), DISTRIBUTIONS, HISTOGRAMS, EMBEDDINGS, and TEXT. On the left side, there is a sidebar with various controls: 'Fit to screen', 'Download PNG', 'Run (1)', 'Session runs (0)', 'Upload' (with a 'Choose File' button), 'Trace Inputs' (a toggle switch), 'Color' (with options for Structure, Device, XLA Cluster, Compute time, and Memory), and 'Colors' (with options for same substructure and unique substructure). At the bottom left, there is a 'Graph' legend with symbols for Namespace, OpNode, Unconnected series, Connected series, Constant, Summary, Dataflow edge, Control dependency edge, and Reference edge. The main area displays a computational graph with two input nodes labeled 'Const' and 'Const_1', each represented by a circle. Arrows from these nodes point to a single output node labeled 'Add', which is represented by a larger oval. The graph is rendered in a light gray color.

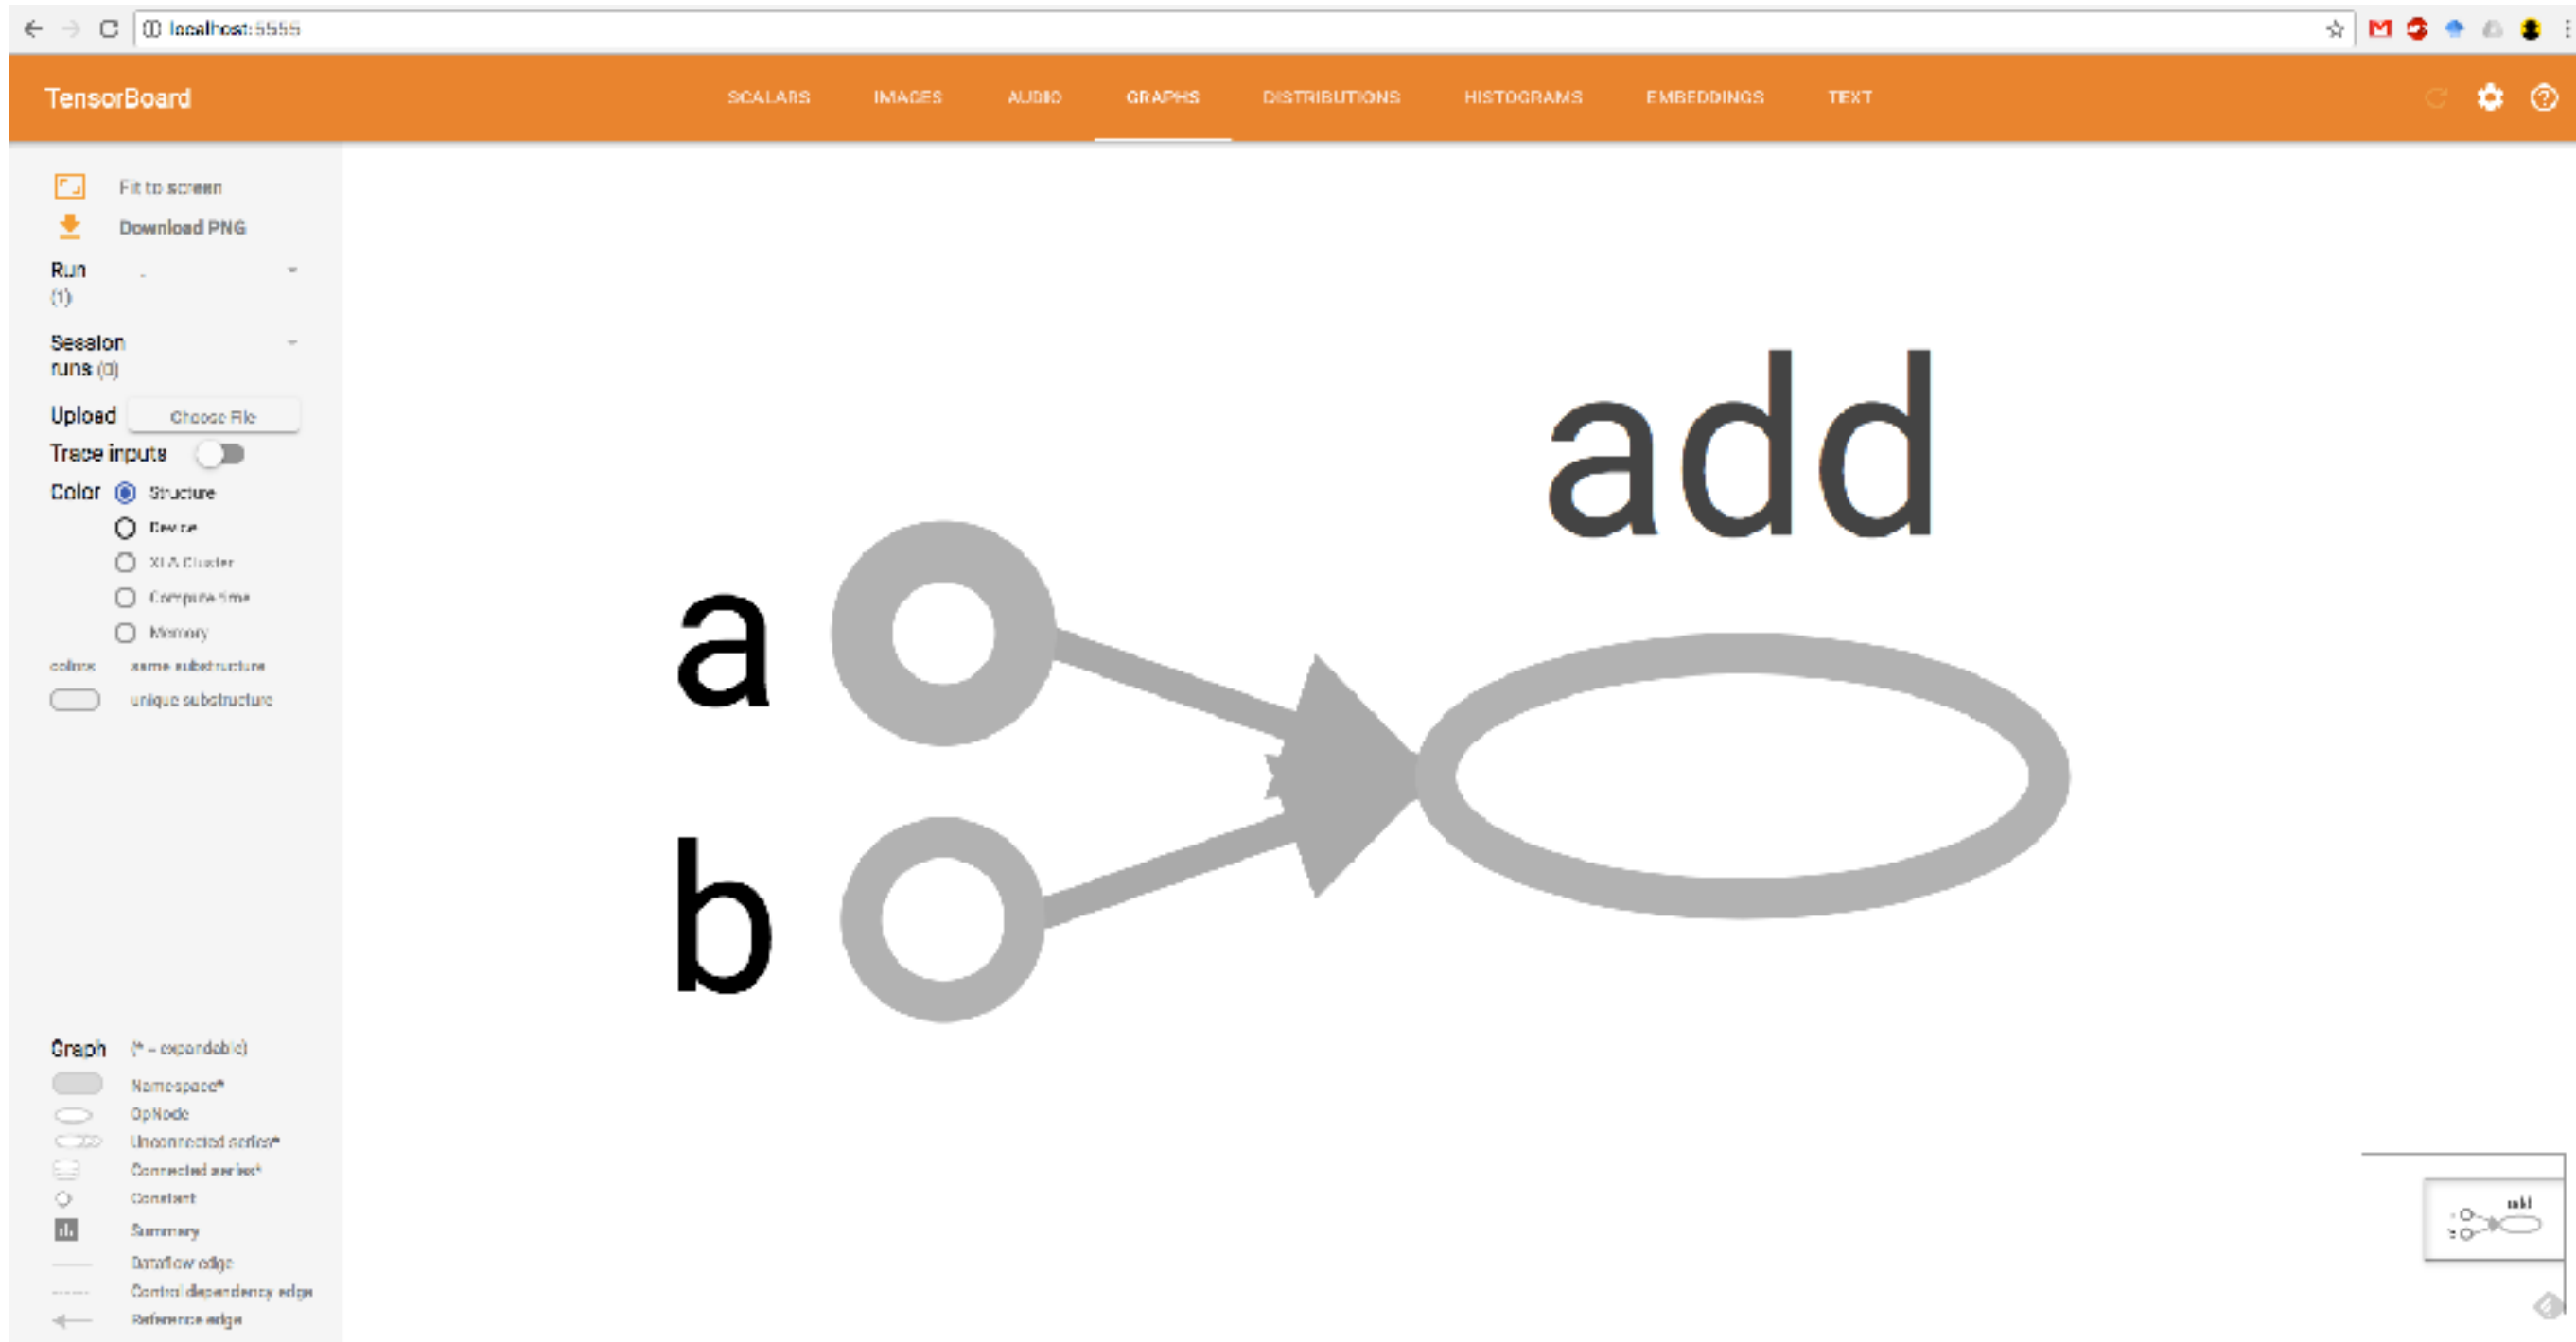
TensorFlow

- TensorBoard for visualization

We can name the variables, as well as the operators:

```
import tensorflow as tf
a = tf.constant(2, name="a")
b = tf.constant(3, name="b")
x = tf.add(a, b, name="add")
with tf.Session() as sess:
    # add this line to use TensorBoard
    writer = tf.summary.FileWriter("./graphs", sess.graph)
    print sess.run(x)
writer.close()
```


TensorFlow



TensorFlow

- Constants

`tf.constant(value, dtype=None, shape=None, name="Const", verify_shape=False)`

- Very similar to that of Numpy
- `tf.zeros`
- `tf.zeros_like`
- `tf.ones`
- `tf.ones_like`
- `tf.fill`
- `tf.constant`
- `tf.linspace`
- `tf.range`

TensorFlow

- Random variables

- `tf.random_normal(shape, mean=0.0, stddev=1.0, dtype=tf.float32, seed=None, name=None)`
- `tf.truncated_normal(shape, mean=0.0, stddev=1.0, dtype=tf.float32, seed=None, name=None)`
- `tf.random_uniform(shape, minval=0, maxval=None, dtype=tf.float32, seed=None, name=None)`
- `tf.random_shuffle(value, seed=None, name=None)`
- `tf.random_crop(value, size, seed=None, name=None)`
- `tf.multinomial(logits, num_samples, seed=None, name=None)`
- `tf.random_gamma(shape, alpha, beta=None, dtype=tf.float32, seed=None, name=None)`
- Set random seed: `tf.set_random_seed(seed)`

https://www.tensorflow.org/api_guides/python/constant_op#Random_Tensors

TensorFlow

- In TensorFlow we can perform all the usual matrix operations in Numpy

Category	Examples
Element-wise mathematical operations	Add, Sub, Mul, Div, Exp, Log, Greater, Less, Equal, ...
Array operations	Concat, Slice, Split, Constant, Rank, Shape, Shuffle, ...
Matrix operations	MatMul, MatrixInverse, MatrixDeterminant, ...
Stateful operations	Variable, Assign, AssignAdd, ...
Neural network building blocks	SoftMax, Sigmoid, ReLU, Convolution2D, MaxPool, ...
Checkpointing operations	Save, Restore
Queue and synchronization operations	Enqueue, Dequeue, MutexAcquire, MutexRelease, ...
Control flow operations	Merge, Switch, Enter, Leave, NextIteration

TensorFlow

- Operations

```
import tensorflow as tf
```

```
a = tf.constant([3, 6])
```

```
b = tf.constant([2, 2])
```

```
tf.add(a, b)    >> [5, 8]
```

```
tf.add_n([a, b, b]) >> [7, 10] = a + b + b
```

```
tf.multiply(a, b) >> [6, 12], elementwise multiplication
```

```
tf.matmul(a, b) >> Error, shape inconsistency for matrix multiplication
```

```
tf.matmul(tf.reshape(a, [1, 2]), tf.reshape(b, [2, 1])) >> 18
```

```
tf.dvi(a, b)    >> [1, 3], elementwise division
```

```
tf.mod(a, b)    >> [1, 0], elementwise modulus
```

- More math operations at:

https://www.tensorflow.org/api_guides/python/math_ops

TensorFlow

- TensorFlow data types

Data type	Python type	Description
<code>DT_FLOAT</code>	<code>tf.float32</code>	32 bits floating point.
<code>DT_DOUBLE</code>	<code>tf.float64</code>	64 bits floating point.
<code>DT_INT8</code>	<code>tf.int8</code>	8 bits signed integer.
<code>DT_INT16</code>	<code>tf.int16</code>	16 bits signed integer.
<code>DT_INT32</code>	<code>tf.int32</code>	32 bits signed integer.
<code>DT_INT64</code>	<code>tf.int64</code>	64 bits signed integer.
<code>DT_UINT8</code>	<code>tf.uint8</code>	8 bits unsigned integer.
<code>DT_UINT16</code>	<code>tf.uint16</code>	16 bits unsigned integer.
<code>DT_STRING</code>	<code>tf.string</code>	Variable length byte arrays. Each element of a Tensor is a byte array.
<code>DT_BOOL</code>	<code>tf.bool</code>	Boolean.
<code>DT_COMPLEX64</code>	<code>tf.complex64</code>	Complex number made of two 32 bits floating points: real and imaginary parts.
<code>DT_COMPLEX128</code>	<code>tf.complex128</code>	Complex number made of two 64 bits floating points: real and imaginary parts.
<code>DT_QINT8</code>	<code>tf.qint8</code>	8 bits signed integer used in quantized Ops.
<code>DT_QINT32</code>	<code>tf.qint32</code>	32 bits signed integer used in quantized Ops.
<code>DT_QUINT8</code>	<code>tf.quint8</code>	8 bits unsigned integer used in quantized Ops.

https://www.tensorflow.org/programmers_guide/dims_types#data_types

TensorFlow

- Variables (tf.Variable is a class, tf. constant is an op)
- Constants are stored in graph definition, but variables are not!

```
# create variable with a scalar value
a = tf.Variable(2, name="scalar")
# create variable with a vector value
b = tf.Variable([2, 3], name="vector")
# create variable with a matrix value
c = tf.Variable([[1, 2], [3, 4]], name="matrix")
# create variable with zeros
W = tf.Variable(tf.zeros([784, 10]))
```

https://www.tensorflow.org/programmers_guide/variables

TensorFlow

- Variables contain operations:

```
x = tf.Variable(...)
```

```
x.initializer # init op
```

```
x.value() # read op
```

```
x.assign(...) # write op
```

```
x.assign_add(...) # and more
```

https://www.tensorflow.org/programmers_guide/variables

TensorFlow

- Variables should be initialized before running

```
# The easiest way to initialize all the variables
```

```
init = tf.global_variables_initializer()
```

```
with tf.Session() as sess:
```

```
    sess.run(init)
```

```
# Initialize a subset of variables
```

```
init_ab = tf.variables_initializer([a, b], name="init_ab")
```

```
with tf.Session() as sess:
```

```
    sess.run(init_ab)
```

```
# Initialize a single variable
```

```
W = tf.Variable(tf.zeros([784, 10]))
```

```
with tf.Session() as sess:
```

```
    sess.run(W.initializer)
```

TensorFlow

- Print the values of variables

```
W = tf.Variable(tf.random_normal([784, 10]))
```

```
with tf.Session() as sess:
```

```
    sess.run(W.initializer)
```

```
    print W
```

```
    print W.eval()
```

- ```
>> <tf.Variable 'Variable_1:0', shape=(784, 10), dtype=float32)
```
- ```
>> [[-0.4778471, -1.3822577, ...]]
```

TensorFlow

- Placeholders
- Symbolic variables that do not take actual value when defined
- Can be filled with actual values when executed
- Examples:

$$f(x, y) = 2x + y$$

- The x , y in the above function are placeholders — they don't have specific values when defined
- However we can evaluate $f(x, y)$ by giving them specific values

TensorFlow

- Placeholders
- **tf.placeholder(dtype, shape=None, name=None)**

```
# create a placeholder of float32, shape is 1x3
```

```
a = tf.placeholder(tf.float32, shape=[3])
```

```
# create a constant of type float32, shape is 1x3
```

```
b = tf.constant([5, 5, 5], tf.float32)
```

```
c = a + b
```

```
with tf.Session() as sess:
```

```
    print sess.run(c)          >> Error?
```

Before running c, we need to provide a with specific values!

TensorFlow

- Placeholders
- **`tf.placeholder(dtype, shape=None, name=None)`**

```
# create a placeholder of float32, shape is 1x3
```

```
a = tf.placeholder(tf.float32, shape=[3])
```

```
# create a constant of type float32, shape is 1x3
```

```
b = tf.constant([5, 5, 5], tf.float32)
```

```
c = a + b
```

```
with tf.Session() as sess:
```

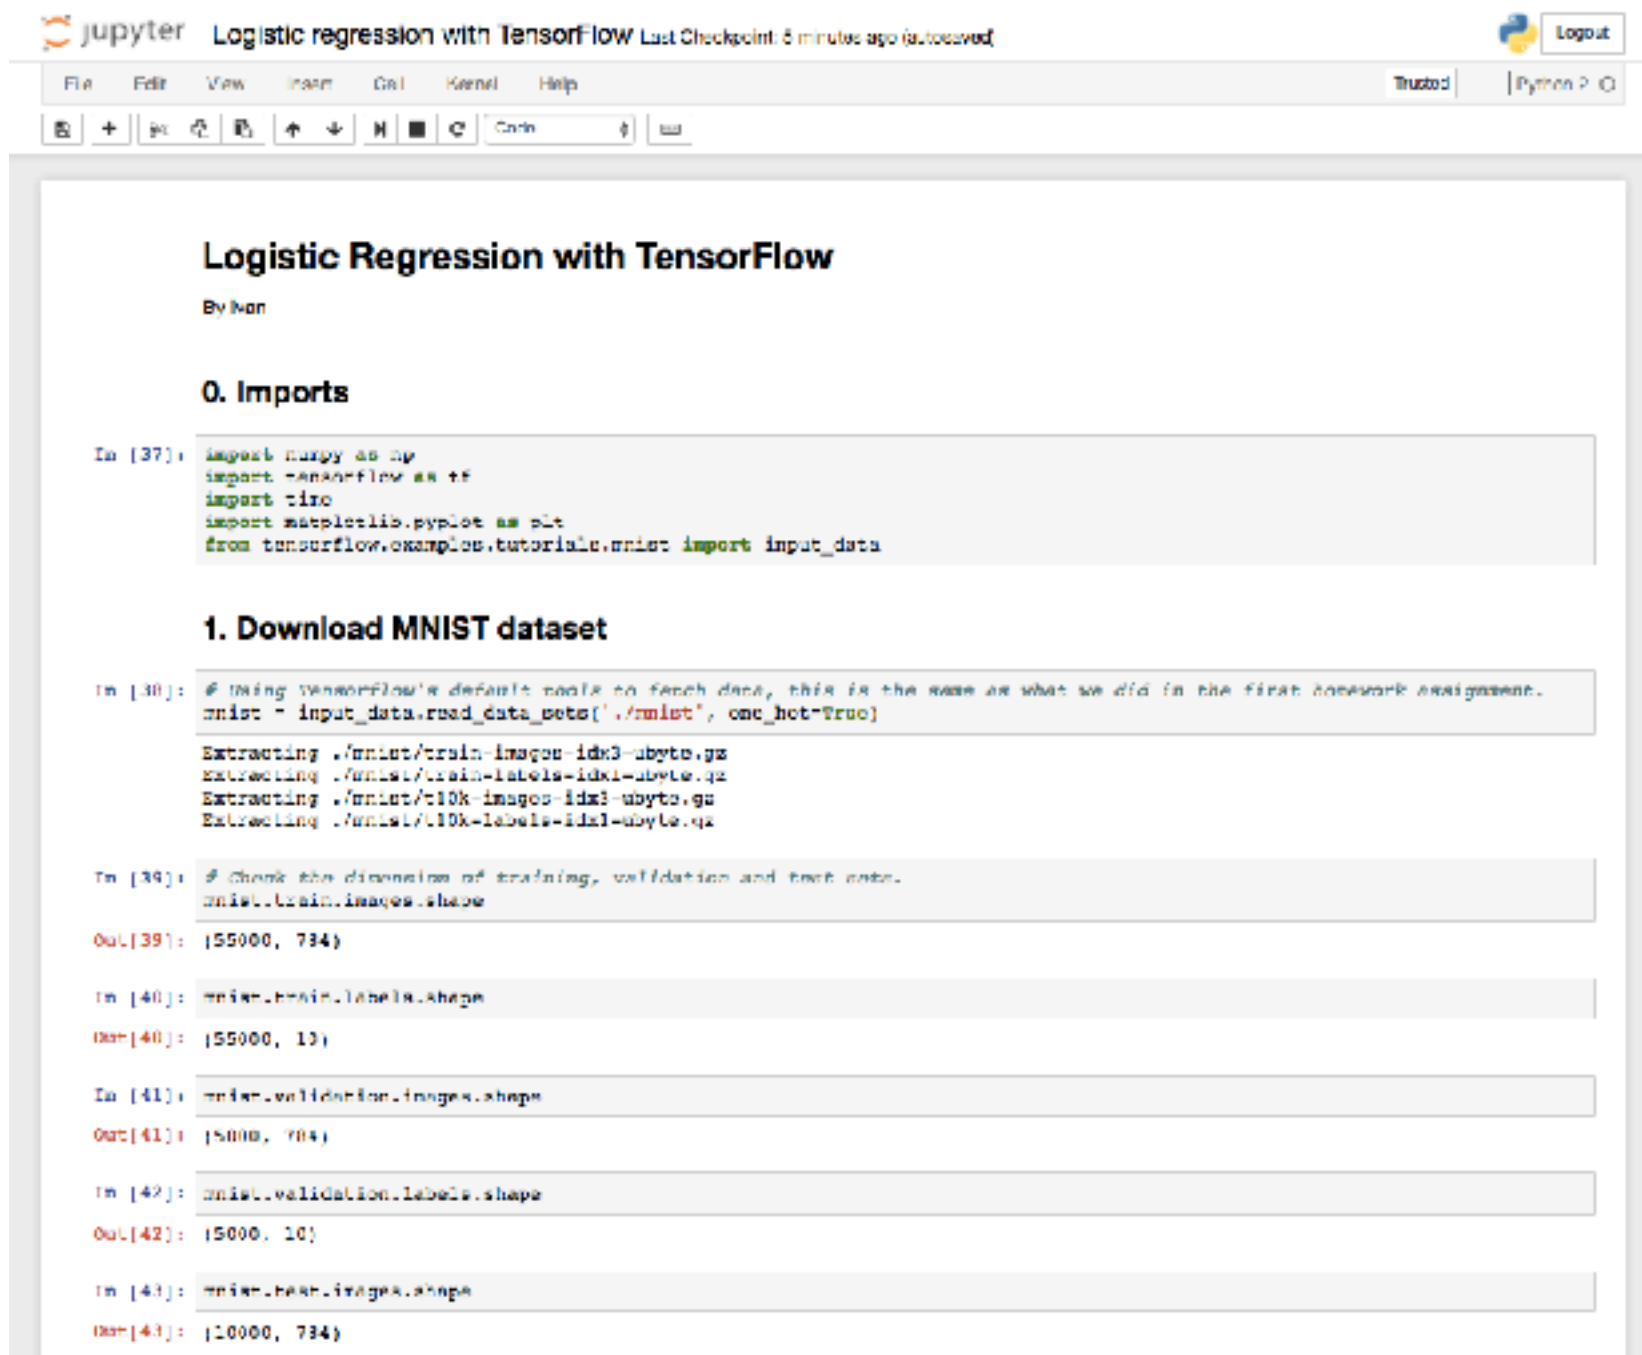
```
    # feed [1, 2, 3] to placeholder a via a dictionary
```

```
    print sess.run(c, {a: [1, 2, 3]}) >> the tensor a can be used as a key, [6, 7, 8]
```

Note: `shape=None` means the placeholder can have any shape

TensorFlow

- Logistic regression



The image shows a Jupyter Notebook interface with the title "Logistic regression with TensorFlow". The notebook is titled "Logistic Regression with TensorFlow" and is by Ivan. It contains the following sections and code:

0. Imports

```
In [37]: import numpy as np
import tensorflow as tf
import time
import matplotlib.pyplot as plt
from tensorflow.examples.tutorials.mnist import input_data
```

1. Download MNIST dataset

```
In [38]: # Using tensorflow's default tools to fetch data, this is the same as what we did in the first homework assignment.
mnist = input_data.read_data_sets('./mnist', one_hot=True)

Extracting ./mnist/train-images-idx3-ubyte.gz
Extracting ./mnist/train-labels-idx1-ubyte.gz
Extracting ./mnist/t10k-images-idx3-ubyte.gz
Extracting ./mnist/t10k-labels-idx1-ubyte.gz

In [39]: # Check the dimensions of training, validation and test sets.
mnist.train.images.shape

Out[39]: (55000, 784)

In [40]: mnist.train.labels.shape

Out[40]: (55000, 10)

In [41]: mnist.validation.images.shape

Out[41]: (5000, 784)

In [42]: mnist.validation.labels.shape

Out[42]: (5000, 10)

In [43]: mnist.test.images.shape

Out[43]: (10000, 784)
```

Homework

- Image classification on MNIST with a three layer MLP:
 - Image size 28x28, size of MLP: 784-500-10, with ReLU as the nonlinear activation function
 - batch size = 200
 - learning rate = 0.1
 - # iterations = 20
 - We provide an initial python script for you to work on
- What you need to implement:
 - Use TensorFlow to build the computational graph
 - Build necessary operator for SGD optimization
 - Run the graph to train the model
 - If implemented correctly, you should see a test set classification accuracy ~ 0.975
 - Running on GTX-1080, taking around ~80 seconds (graph compilation take time).
- You need to install numpy and tensorflow



扫码获取最新数据科学资讯